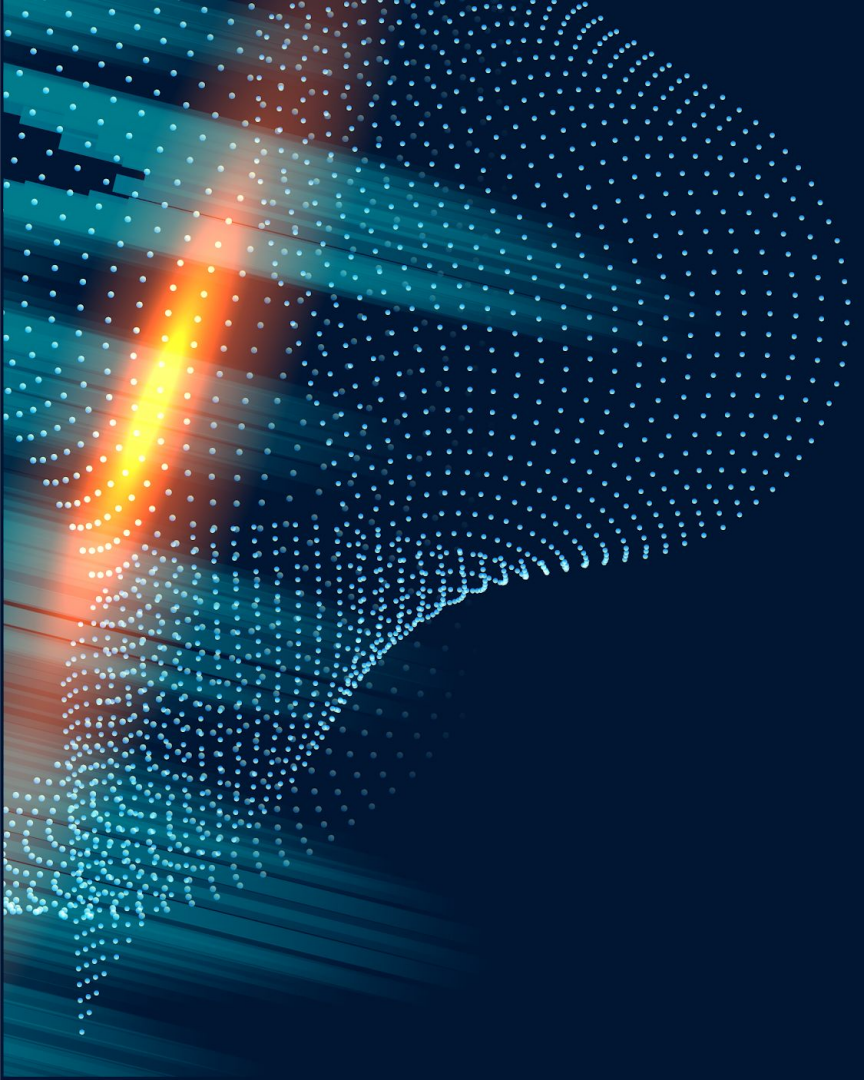


IMAGE PROCESSING MINI PROJECT

EC386

PM Prasanna 201EC242
Abhinav Raghunandan 201EC102

An abstract graphic on the left side of the slide, featuring a series of concentric, semi-circular arcs composed of small white dots. A bright, glowing orange and yellow light source is positioned on the left, casting a beam of light across the arcs. The background is a dark blue gradient.

CAPSULE ENDOSCOPY

Image Classification Techniques



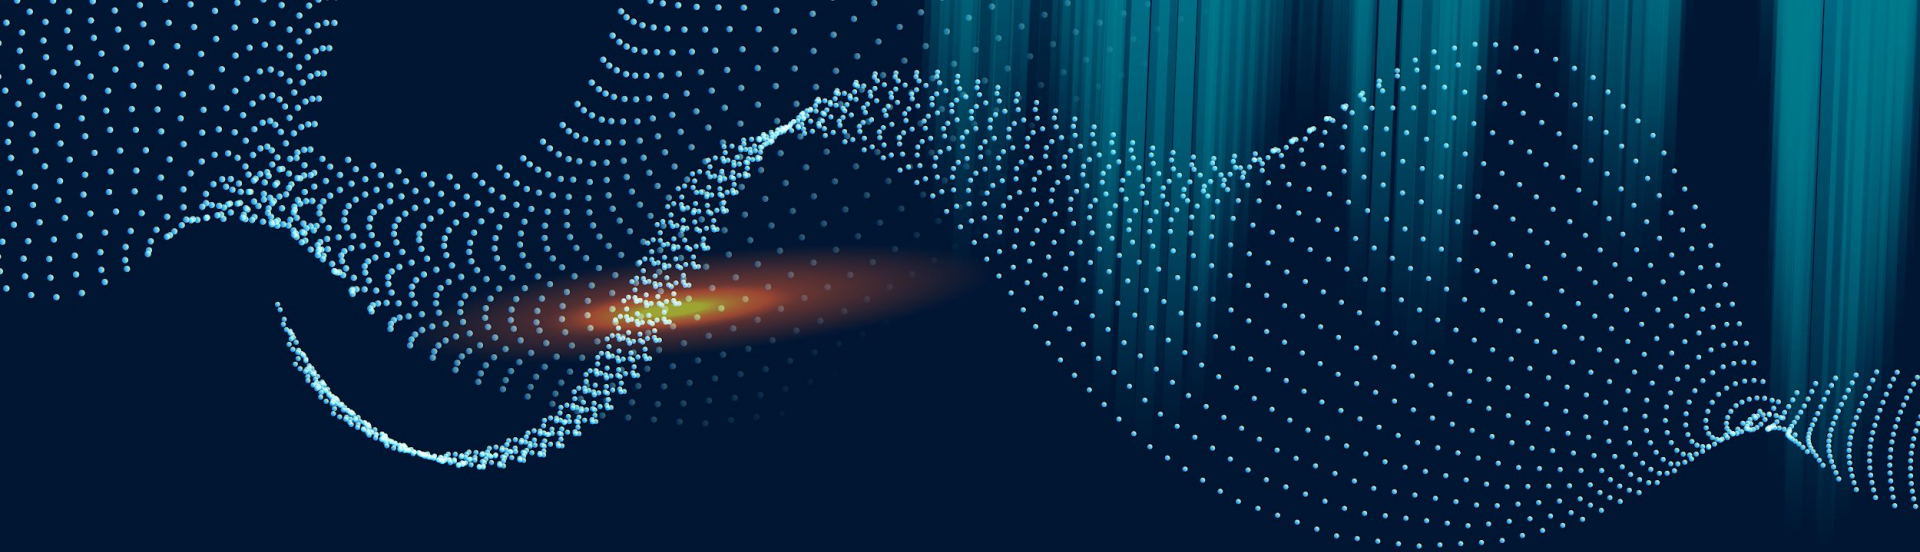
01

PROBLEM STATEMENT

Brief description about the Aim
of the this project.

**Perform Anatomical Classification
using interior images of the intestine
obtained from the Capsule Endoscopy
Technique**



**The intestine contains 3 major regions
- Pylorus, Cecum and the Z- Line.
Perform Image Classification based on
texture, size and other features which
is useful for medical treatments.**



02

DATASET

- The dataset contains images of the different sections of the colon captured using Capsule Endoscopy
- Source of Dataset : Dr. Aparna P Dinesh (Project Mentor)



Data Augmentation

- Rotation, Zooming in and out, shifting height and width and horizontal flipping of the image
- Performed Data augmentation and finally obtained 131 images belonging to 3 classes as 'Training' data
- And 31 images belonging to 3 classes as Validation data.





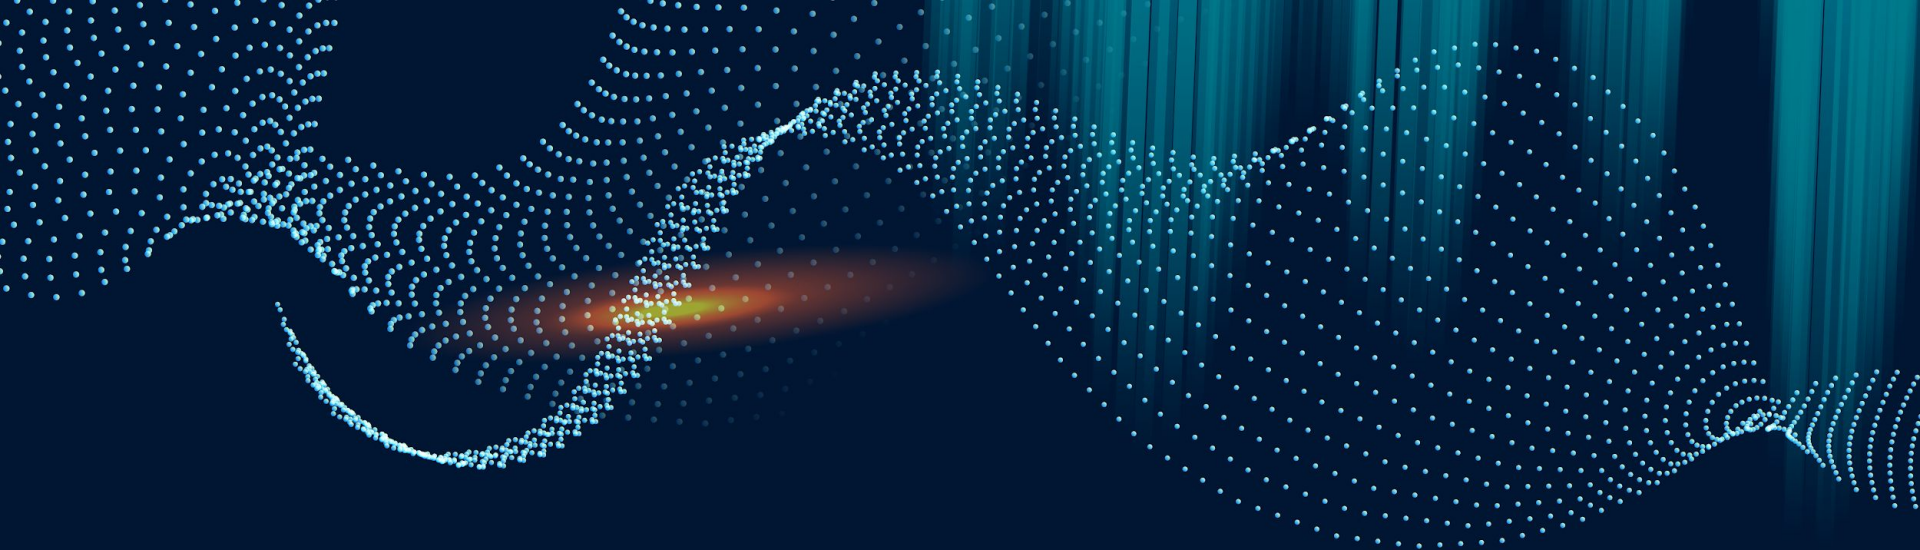
```
# create a data generator
```

```
datagen = ImageDataGenerator(  
    samplewise_center=True, # set each sample mean to 0  
    rotation_range=10, # randomly rotate images in the range (degrees, 0 to 180)  
    zoom_range = 0.2, # Randomly zoom image  
    width_shift_range=0.2, # randomly shift images horizontally (fraction of total width)  
    height_shift_range=0.2, # randomly shift images vertically (fraction of total height)  
    horizontal_flip=True, # randomly flip images  
    vertical_flip=False)
```

```
[ ] train_it = datagen.flow_from_directory('/content/gdrive/MyDrive/Mini_Project/Train',  
                                          target_size=(224, 224),  
                                          color_mode='rgb',  
                                          class_mode='categorical',  
                                          batch_size=5)
```

```
# load and iterate validation dataset
```

```
valid_it = datagen.flow_from_directory('/content/gdrive/MyDrive/Mini_Project/Test',  
                                       target_size=(224, 224),  
                                       color_mode='rgb',  
                                       class_mode='categorical',  
                                       batch_size=5)
```

03

**PREVIOUS
WORK**

Transfer Learning Model Implementation using VGG- 16

- Classified the image dataset using VGG-16 model.


Reasons for choosing VGG-16


- Small dataset used
- Highly efficient and accurate as it is a pre-trained model

[GitHub Repo Code Link VGG16](#)



Accuracy Graph Plot

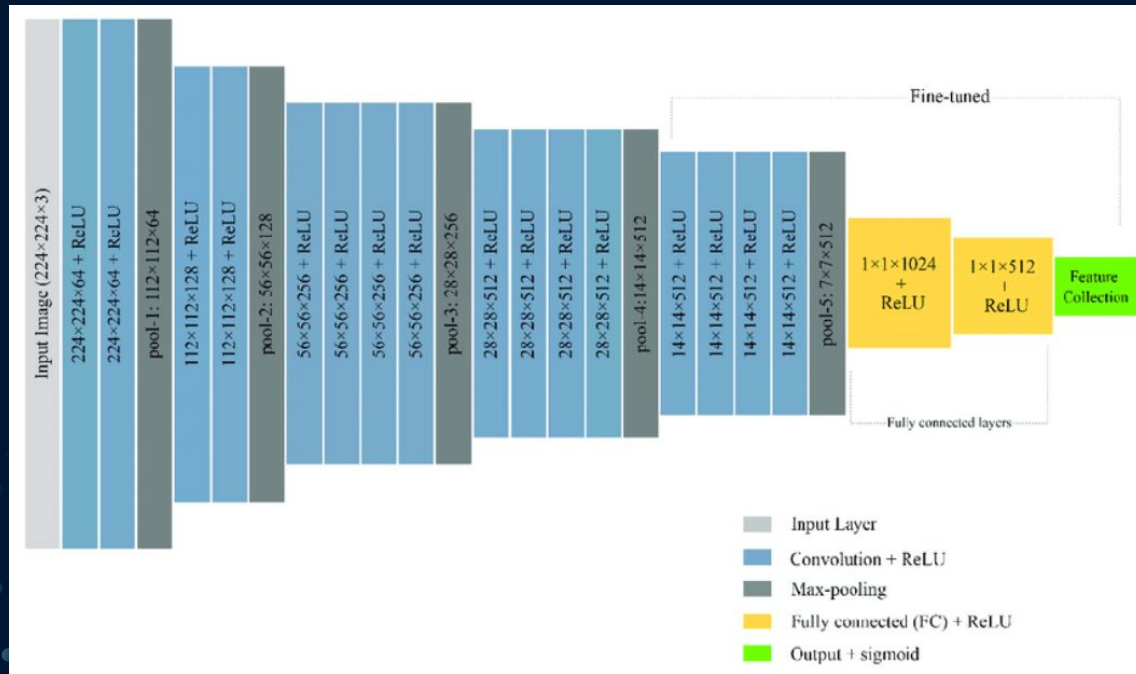




04

**RECENT
WORK**

VGG-19 MODEL



Using VGG- 19 Model

- Classified the image dataset using VGG-19 model.
- Model was suggested by our Project Mentor

Advantages of VGG-19

- 3 more Convolutional Layers than the VGG-16 Model
- Bound to result in Higher Accuracy

[GitHub Repo Code Link VGG19](#)



VGG-19 Model Summary

Model: "sequential"

Layer (type)	Output Shape	Param #
vgg19 (Functional)	(None, 7, 7, 512)	20024384
flatten (Flatten)	(None, 25088)	0
dense (Dense)	(None, 512)	12845568
dense_1 (Dense)	(None, 128)	65664
dense_2 (Dense)	(None, 3)	387

=====
Total params: 32,936,003


Trainable params: 12,911,619

Non-trainable params: 20,024,384
=====

Accuracy Graph Plot

2s 226ms/step - loss: 1.3921 - accuracy: 0.9200 - val_loss: 4.9863 - val_accuracy: 0.8667

3s 308ms/step - loss: 0.7570 - accuracy: 0.9362 - val_loss: 4.5534 - val_accuracy: 0.6667



Model Implementation

```
▶ VGG19_model = Sequential()  
  pretrained_model = keras.applications.vgg19(  
    weights='imagenet', # Load weights pre-trained on ImageNet.  
    input_shape=(224, 224, 3),  
    include_top=False)
```

```
⬤ Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/vgg19/80134624/80134624 [=====] - 4s 0us/step
```

```
[ ] for layer in pretrained_model.layers:  
    layer.trainable=False
```

```
VGG19_model.add(pretrained_model)
```

```
[ ] VGG19_model.add(Flatten())  
    VGG19_model.add(Dense(512, activation='relu'))  
    VGG19_model.add(Dense(128, activation='relu'))  
    VGG19_model.add(Dense(3, activation='softmax'))
```

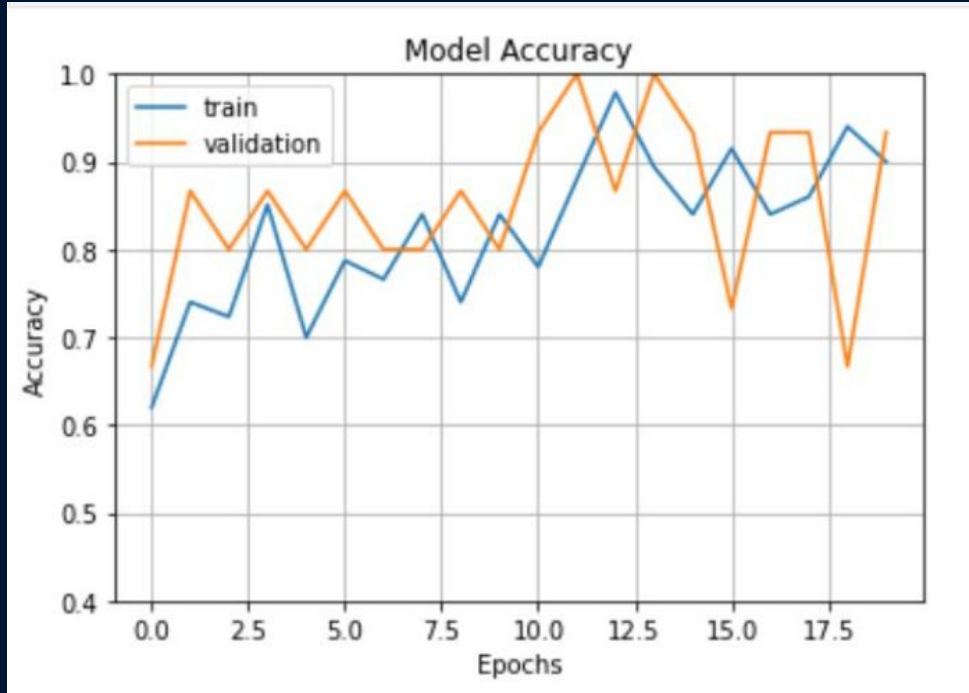
EfficientNet B0

EfficientNet B0

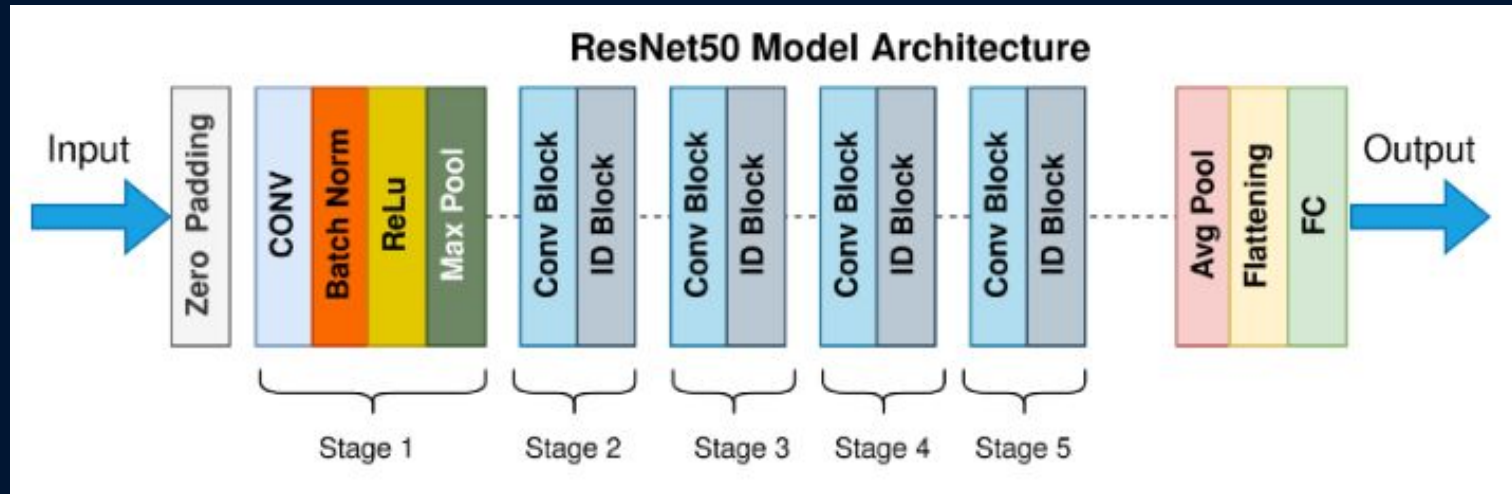
- [Paper published on EfficientNet and ResNet Networks](#)
- Lower accuracy than expected since we run the risk of Overfitting when we use EfficientNet with small dataset applications

[GitHub Repo Code Link EfficientNet](#)

Code and Implementation



ResNet-50 MODEL



Using ResNet-50 Model

- ResNet-50 is a Convolutional Neural Network which is 50 layers deep
- Pre- trained model has trained over a million images using the ImageNet Database

Reasons for choosing ResNet-50

Article about ResNet enhancing performance

- Significant enhancement of performance due to more convolutional layers
- Maintain low error rate in the deep layers of the network

[GitHub Repo Code Link ResNet50](#)




ResNet-50 Model Summary

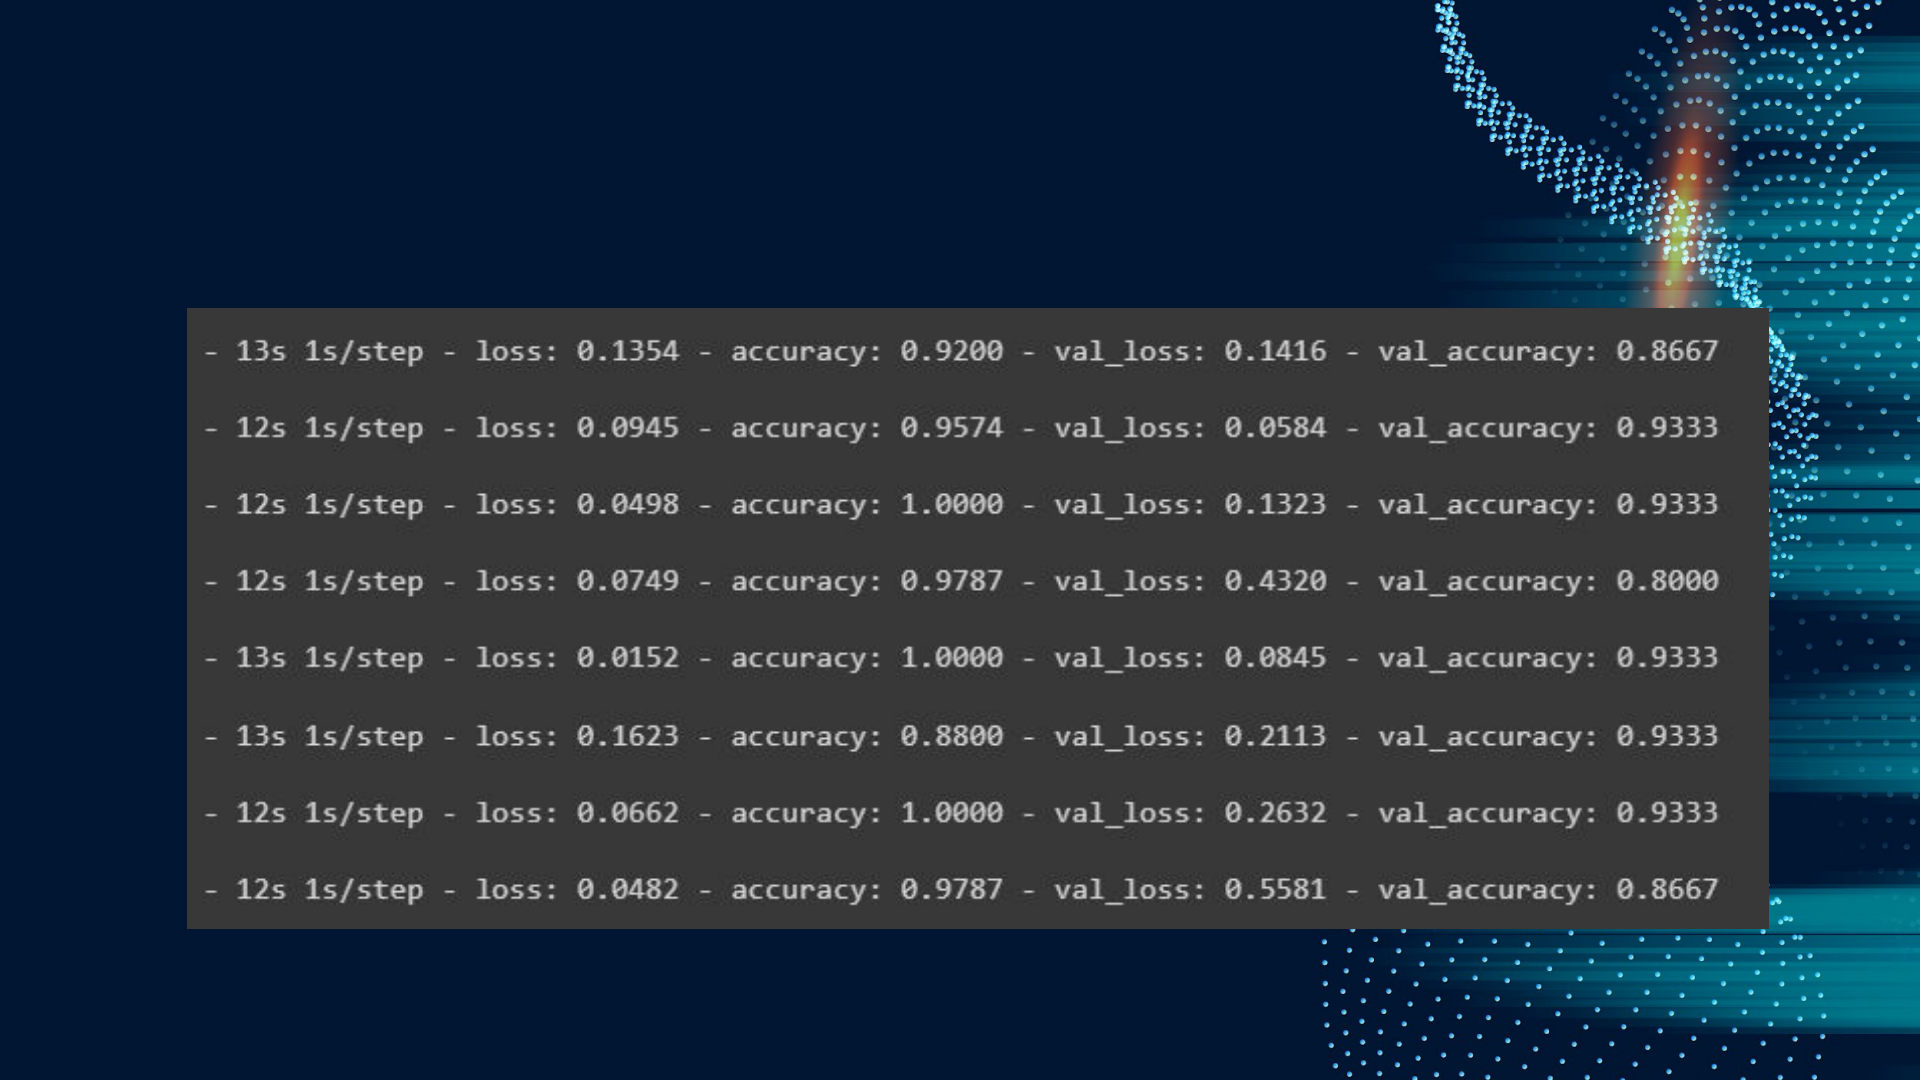
Model: "sequential"

Layer (type)	Output Shape	Param #
resnet50 (Functional)	(None, 2048)	23587712
flatten (Flatten)	(None, 2048)	0
dense (Dense)	(None, 512)	1049088
dense_1 (Dense)	(None, 128)	65664
dense_2 (Dense)	(None, 3)	387

=====
Total params: 24,702,851
Trainable params: 1,115,139
Non-trainable params: 23,587,712

Accuracy Graph Plot





```
- 13s 1s/step - loss: 0.1354 - accuracy: 0.9200 - val_loss: 0.1416 - val_accuracy: 0.8667
- 12s 1s/step - loss: 0.0945 - accuracy: 0.9574 - val_loss: 0.0584 - val_accuracy: 0.9333
- 12s 1s/step - loss: 0.0498 - accuracy: 1.0000 - val_loss: 0.1323 - val_accuracy: 0.9333
- 12s 1s/step - loss: 0.0749 - accuracy: 0.9787 - val_loss: 0.4320 - val_accuracy: 0.8000
- 13s 1s/step - loss: 0.0152 - accuracy: 1.0000 - val_loss: 0.0845 - val_accuracy: 0.9333
- 13s 1s/step - loss: 0.1623 - accuracy: 0.8800 - val_loss: 0.2113 - val_accuracy: 0.9333
- 12s 1s/step - loss: 0.0662 - accuracy: 1.0000 - val_loss: 0.2632 - val_accuracy: 0.9333
- 12s 1s/step - loss: 0.0482 - accuracy: 0.9787 - val_loss: 0.5581 - val_accuracy: 0.8667
```


Model Implementation

conv5_block1_2_relu (Activation)	(None, 7, 7, 512)	0	['conv5_block1_2_bn[0][0]']
conv5_block1_0_conv (Conv2D)	(None, 7, 7, 2048)	2099200	['conv4_block6_out[0][0]']
conv5_block1_3_conv (Conv2D)	(None, 7, 7, 2048)	1050624	['conv5_block1_2_relu[0][0]']

```
for layer in pretrained_model.layers:  
    layer.trainable=False
```

```
resnet_model.add(pretrained_model)
```

```
resnet_model.add(Flatten())  
resnet_model.add(Dense(512, activation='relu'))  
resnet_model.add(Dense(128, activation='relu'))  
resnet_model.add(Dense(3, activation='softmax'))
```

Confusion Matrix for the ResNet-50 Model

	CECUM	PYLORUS	Z- LINE
cecum1.png	0.998	1.65×10^{-8}	5.22×10^{-14}
pylorus4.png	7.07×10^{-9}	0.999	10^{-12}
z_line7.png	8.91×10^{-14}	2.08×10^{-8}	0.999

```

print(v_pred)
[[9.98878539e-01 1.11226016e-03 9.18452406e-06]
[9.96178746e-01 3.82050802e-03 6.88982595e-07]
[9.99994755e-01 5.24385132e-06 1.02709219e-08]
[9.97927189e-01 1.32980838e-03 7.43058743e-04]
[9.99959469e-01 4.04923485e-05 5.95777934e-08]
[9.98265445e-01 1.72999338e-03 4.65316407e-06]
[9.99788582e-01 2.11322636e-04 1.32466411e-07]
[9.99925375e-01 7.45900834e-05 5.59755335e-08]
[9.99896049e-01 9.56724107e-05 8.25873758e-06]
[9.99989748e-01 1.01846617e-05 8.39739158e-08]
[1.68626080e-04 9.97170985e-01 2.66038673e-03]
[6.85664418e-05 9.46613193e-01 5.33181429e-02]
[2.14870670e-03 9.96903241e-01 9.48033412e-04]
[9.02599422e-04 9.98030007e-01 1.06749509e-03]
[9.82837635e-04 9.98999417e-01 1.77204747e-05]
[1.60931580e-04 9.99523401e-01 3.15610232e-04]
[2.69379176e-04 9.99715984e-01 1.46796665e-05]
[5.59857071e-05 9.99632597e-01 3.11383279e-04]
[3.88873072e-04 9.99452889e-01 1.58249313e-04]
[1.02829697e-04 9.82921839e-01 1.69752371e-02]
[1.67142425e-04 2.56699808e-02 9.74162877e-01]
[4.97193575e-09 2.60668003e-05 9.99973893e-01]
[2.80589175e-05 2.52819378e-02 9.74690020e-01]
[1.16978108e-05 2.27108225e-01 7.72880018e-01]
[1.75953119e-05 3.91543144e-03 9.96066988e-01]
[1.75027594e-07 9.08607326e-04 9.99091268e-01]
[5.66822791e-06 3.11501473e-02 9.68844235e-01]
[7.92791681e-08 1.57806743e-03 9.98421907e-01]
[5.66583580e-08 9.96078888e-05 9.99900341e-01]
[6.59082167e-09 1.01506732e-04 9.99898434e-01]

```

11 22

\rightarrow χ^2 cum
 \rightarrow Pythorus
 \rightarrow Z-Line

$\left. \begin{array}{l} \text{---} \\ \text{---} \\ \text{---} \end{array} \right\} = 22 \ 1$



98.84%

Accuracy of our ResNet- 50
Model

SUMMARY OF MODELS



VGG-16

Accuracy
88.45%



VGG-19

Accuracy
93.62%



EfficientNet
B0

Accuracy
95.79%



ResNet-50

Accuracy
98.84%

REFERENCES

1. [Introduction to Image Classification](#)
2. [About Data Augmentation](#)
3. [About Activation Functions](#)
4. [Transfer Learning](#)
5. [VGG-19 Model](#)
6. [Paper on EfficientNet and ResNets](#)
7. [Parameters in Convolutional Neural Networks](#)
8. [StackOverflow for General Error Handling](#)





THANK YOU

PM Prasanna 201EC242
Abhinav Raghunandan 201EC102